

DATA MINING PROJECT



JANUARY 1, 2020

AIN SHAMS UNIVERSITY – FACULTY OF ENGINEERING

FACULTY OF ENGINEERING

AIN SHAMS UNIVERSITY

CSE 385: Data Mining & Business Intelligence



Data Mining Project

Submitted by:

Abdelrahman Amr Issawi

aid-issawi@hotmail.com

16P6001@eng.asu.edu.eg

Ahmed Hesham AlSaey

ahmed.alsai@hotmail.com

16P6000@eng.asu.edu.eg

Sherif Ashraf Ahmed

sherifashraf_98@hotmail.com

16P9033@eng.asu.edu.eg

Mayar Khaled Mohamed

MayarKhaledd_3@yahoo.com

16P6030@eng.asu.edu.eg

Submitted to:

Dr. Mahmoud Mounir,

Eng. Ahmed.

JANUARY 1, 2020

A report for Data Mining & Business Intelligence Course coded CSE385 with the requirements of Ain Shams University.

Table of Contents

1.0 INTRODUCTION.....	3
2.0 BRIEF DESCRIPTION.....	4
Implementation	4
Added Features “Bonus”	5
3.0 USER GUIDE.....	7
4.0 CODES	8
Main.py	8
Preprocessing.py.....	12
regression.py.....	14
classification.py.....	15
clustering.py.....	15

1.0 INTRODUCTION

In this project we are applying the algorithms we studied. Regression, classification and clustering. Also, we are pre-processing the data before we apply any of the algorithms, and finally we get the results of applying these tools.

Regression and classification are supervised learning approaches that map an input to an output based on example input-output pairs, while clustering is an unsupervised learning approach.

Regression: It predicts continuous valued output. Regression analysis is the statistical model which is used to predict the numeric data instead of labels. It can also identify the distribution trends based on the available data or historic data. Predicting a person's income from their age, education is example of regression task.

Classification: It predicts discrete number of values. In classification the data is categorized under different labels according to some parameters and then the labels are predicted for the data. Classifying emails as either spam or not spam is example of classification problem.

Clustering: Clustering is the task of partitioning the dataset into groups, called clusters. The goal is to split up the data in such a way that data sets within single cluster are very similar and data sets in different clusters are different. It determines grouping among unlabeled data.

Classification is applied on the cancer dataset, Regression is applied on diamonds dataset, and Clustering is applied on the iris dataset.

In the report attached we are discussing our implementation, the added features, you will also find a user guide, and some test cases. At the end you will find our codes.

Anyway, I hope this quick introduction has helped you, Now let's read the report.

2.0 BRIEF DESCRIPTION

Implementation

There are 5 python files Main.py, Preprocessing.py, regression.py, classification.py, clustering.py. Main.py is the main function calling the other functions. The main function reads the data files needed for classification, regression, and clustering as well.

For classification and regression, the main function calls the **dataCleaning()** function defined in pre-processing class it splits the data set into training and testing data, and drops the unneeded columns. Then, the 2 classes are sent to the **encode()** function giving values to any un numerical data. Then the missing values dealing method is specified, whether to **drop()** the tuple, **replaceMean()**, or **replaceMode()**. The data set is **split()** into X_train_class, X_test_class, y_train_class, y_test_class. And finally the data is scaled using **scale()** using StandardScaler, MinMaxScaler, MaxScaler, or RobustScaler.

Then the data is ready for classification or regression, for classification the function **classify()** is called taking the chosen algorithm to classify with, X_train_class, y_train_class to train the machine, then we classify the X_test_class and gives the results to **calculateAccuracy()** calculateAccuracy to compare it with the y_test_class.

The same steps are done in the regression as well but we use the function **predict()**, and **getScore()**.

For clustering, we load the dataset to X_clust, y_clust, then we use the **splitCluster()** function to break the dataset into training and testing data. We **scale()** the data, and finally we **predict()** the clusters.

Implemented classifying methods:

- KNN,
- Decision Tree,
- Naive Bayes,
- Random Forest.

Implemented regression methods:

- Linear Regression,
- Decision Tree,
- Polynomial Regression,
- KNN Regression,
- Random Forest.

Implemented clustering methods:

- K-means.

Added Features “Bonus”

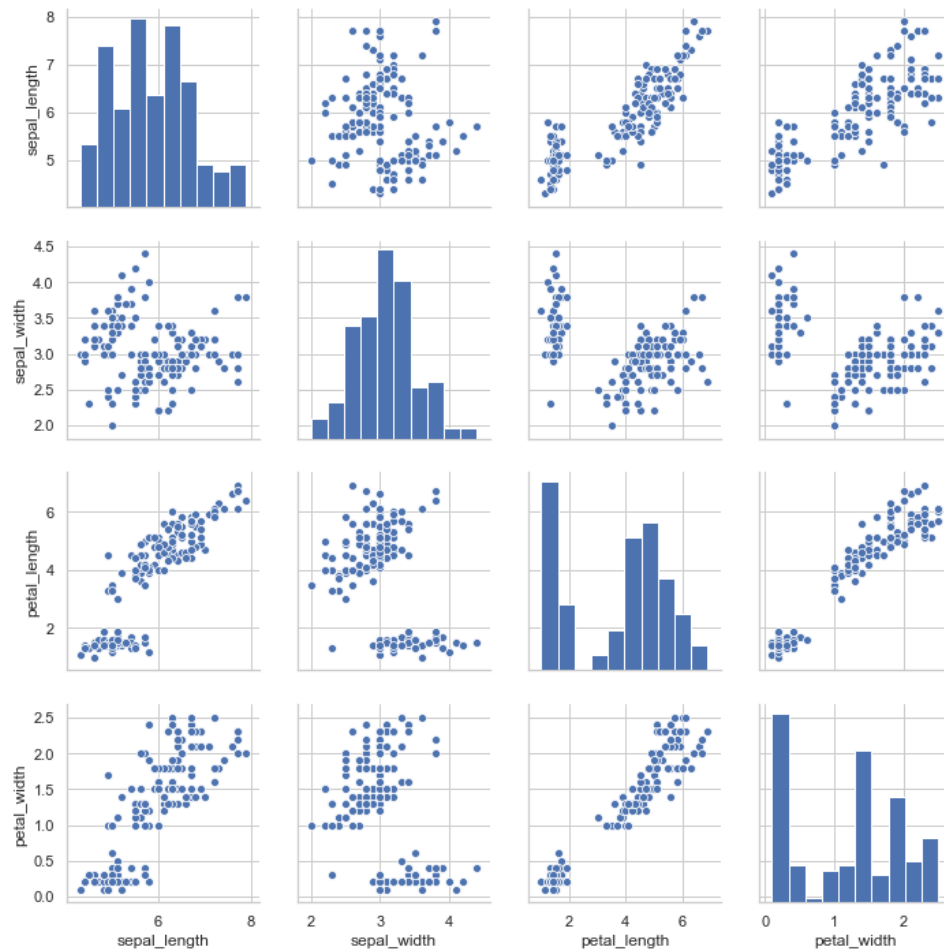


- Extra regression algorithm is added “Random Forest”
- GUI is implemented so that it’s easier to choose the algorithm you need to apply, also you can write multiple missing values dealing methods, and multiple scaling methods. The score, accuracy, and clusters are displayed in it. The GUI is implemented using tkinter library in python, it’s defined to use the labels, buttons, entry fields, and list-box.

The screenshot shows a window titled "DM" with a light gray background. It contains three rows of configuration options, each with a label, a dropdown menu, and a text entry field. The first row is for "Classifier Method" with a dropdown set to "KNN" and a "Missing Values" field set to "drop". The second row is for "Regression Method" with a dropdown set to "Linear Regression" and a "Missing Values" field set to "drop". The third row is for "Clustering Method" with a dropdown set to "K-Means" and an empty "Missing Values" field. Each row also has a "Scaling Method" dropdown set to "StandardScaler". A large blue "Submit" button is centered below the configuration options. Below the button, there are two lines of text: "***Scaling Methods are: StandardScaler, MinMaxScaler, or MaxScaler" and "***We deal with missing values by: drop, replaceMean, or replaceMode". At the bottom, there is a large white text area for output.

- High classifying accuracy.

This screenshot shows the same "DM" window as the previous one, but with the results of the KNN algorithm displayed in the output text area. The configuration remains the same: Classifier Method is KNN, Regression Method is Linear Regression, and Clustering Method is K-Means. The output text area contains the following text: "***Scaling Methods are: StandardScaler, MinMaxScaler, or MaxScaler", "***We deal with missing values by: drop, replaceMean, or replaceMode", "The accuracy is:99.41520467836257% using KNN Algorithm", "The score is: 0.8737775155528267 using Linear Regression Algorithm", and "Predicted clusters are [1 2 0 2 0 2 0 0 1 1 1 0 2 1 2 2 0 0 2 2 0 2 0 2 1 2 0 0 1 1 2 0 0 0 1 0 2 0 1 0 1 2] using K-Means Algorithm". The output text area has a vertical scrollbar on the right side.

- Clustering visualization.



3.0 USER GUIDE

DM

Classifier Method: Random Forest Scaling Method: StandardScaler Missing Values: replaceMean

Regression Method: Linear Regression Scaling Method: MinMaxScaler Missing Values: replaceMode

Clustering Method:
Linear Regression
Decision Tree
Polynomial Regression
KNN Regression
Random Forest

Submit

***Scaling Methods are: StandardScaler, MinMaxScaler, or MaxScaler
***We deal with missing values by: drop, replaceMean, or replaceMode

The accuracy is: 97.07602339181285% using Random Forest Algorithm

The score is: 0.873777515552827 using Linear Regression Algorithm

Predicted clusters are [0 2 1 2 1 2 1 2 0 1 1 1 0 2 2 2 0 0 0 1 1 1 1 0 2 0 1 1 0 0 0 1 0 2 1 2 0 2 2 1 1 0 2] using K-Means Algorithm

1. Choose the Algorithm.
2. Choose the scaling method.
3. Choose the missing values dealing method.
4. Click submit.
5. Observe your results.

4.0 CODES

Main.py

```
import pandas as pd
import numpy as np
from preprocessing import Preprocessor
from classification import Classifier
from regression import Regression
from clustering import Clustering
from sklearn.datasets import load_iris
import seaborn as sns
import matplotlib.pyplot as plt
from tkinter import *

def buttonClick():
    classificationAlg = variable1.get()
    regressionAlg = variable2.get()
    clusterAlg = variable3.get()

    classificationScaler = box1s.get()
    regressionScaler = box2s.get()
    clusterScaler = box3s.get()

    classificationMethod = box1F.get()
    regressionMethod = box2F.get()

    preprocessing = Preprocessor()

    # Classification
    classificationData = pd.read_csv('cancer.csv')
    X_class, y_class = preprocessing.dataCleaning(classificationData, 'cancer')
    X_class, y_class = preprocessing.encode(X_class, y_class, 'cancer')
    if classificationMethod == "drop":
        classificationData = preprocessing.drop(classificationData)
    elif classificationMethod == "replaceMean":
        classificationData = preprocessing.replaceMean(classificationData)
    elif classificationMethod == "replaceMode":
        classificationData = preprocessing.replaceMode(classificationData)
    else:
        classificationData = preprocessing.drop(classificationData)
    X_train_class, X_test_class, y_train_class, y_test_class =
preprocessing.split(X_class, y_class, 0.3)
    X_train_class, X_test_class = preprocessing.scale(X_train_class, X_test_class,
type=classificationScaler)
    classifier = Classifier(classificationAlg, X_train_class, np.ravel(y_train_class))
    y_pred_class = classifier.classify(X_test_class)
    accuracy = classifier.calculateAccuracy(np.ravel(y_test_class), y_pred_class)

    # Regression
    regressionData = pd.read_csv('diamonds.csv')
    X_regression, y_regression = preprocessing.dataCleaning(regressionData,
'diamonds')
    X_regression, y_regression = preprocessing.encode(X_regression, y_regression,
'diamonds')
    if regressionMethod == "drop":
        regressionData = preprocessing.drop(regressionData)
    elif regressionMethod == "replaceMean":
        regressionData = preprocessing.replaceMean(regressionData)
    elif regressionMethod == "replaceMode":
        regressionData = preprocessing.replaceMode(regressionData)
    else:
```

```

        regressionData = preprocessing.drop(classificationData)
        X_train_reg, X_test_reg, y_train_reg, y_test_reg =
preprocessing.split(X_regression, y_regression, 0.3)
        X_train_reg, X_test_reg = preprocessing.scale(X_train_reg, X_test_reg,
type=regressionScaler)
        regression = Regression(regressionAlg, X_train_reg, np.ravel(y_train_reg))
        y_pred_regression = regression.predict(X_test_reg)
        score = regression.getScore(np.ravel(y_test_reg), y_pred_regression)

        # Clustering
        X_clust, y_clust = load_iris(return_X_y=True)
        iris = sns.load_dataset("iris")
        X_train_clust, X_test_clust = preprocessing.splitCluster(X_clust, 0.3)
        X_train_clust, X_test_clust = preprocessing.scale(X_train_clust, X_test_clust,
type=clusterScaler)
        cluster = Clustering(clusterAlg, X_train_clust)
        y_pred_cluster = cluster.predict(X_test_clust)

        print('\n\nThe accuracy is: ' + str(accuracy) + ' % ' + str(classificationAlg))
        print('\n\nThe score: ' + str(score) + ' ' + str(regressionAlg))
        print('\n\nPredicted clusters: ' + str(y_pred_cluster) + ' ' + str(clusterAlg))

        content = "The accuracy is:" + str(accuracy) + "%    using " + str(
            classificationAlg) + " Algorithm \n\n" + "The score is: " + str(score) + "
using " + str(
            regressionAlg) + " Algorithm \n\n" + "Predicted clusters are " +
str(y_pred_cluster) + " using " + str(
            clusterAlg) + " Algorithm \n\n"
        text1 = Text(window, width=80, height=8)
        text1.grid(row=9, column=0, columnspan=30)
        text1.insert(INSERT, content)
        sb1 = Scrollbar(window)
        sb1.grid(row=9, column=5, columnspan=40)
        text1.configure(yscrollcommand=sb1.set)
        sb1.configure(command=text1.yview)
        window.update()
        # Construct iris plot
        sns.swarmplot(x="species", y="petal_length", data=iris)
        # Show plot
        plt.show()

        sns.set(style="whitegrid")
        sns.boxplot(data=iris, orient="h", palette="Set2")
        plt.show()

        sns.set(style="whitegrid")
        sns.pairplot(iris)
        plt.show()

window = Tk()
window.geometry("877x320")
window.title("DM")
window.resizable(False, False)

variable1 = StringVar(window)
variable1.set("KNN") # default value
l1 = Label(window, text="Classifier Method: ", width=20)
l1.grid(row=0, column=0)
box1 = OptionMenu(window, variable1, "KNN", "Decision Tree", "Naive Bayes", "Random
Forest")
box1.grid(row=0, column=1)
l1s = Label(window, text="Scaling Method: ", width=20)

```

```

l1s.grid(row=0, column=2)
box1s = Entry(window)
box1s.grid(row=0, column=3)
box1s.insert(END, "StandardScaler")

l1F = Label(window, text="Missing Values: ", width=20)
l1F.grid(row=0, column=4)
box1F = Entry(window)
box1F.insert(END, "drop")
box1F.grid(row=0, column=5)

variable2 = StringVar(window)
variable2.set("Linear Regression") # default value
l2 = Label(window, text="Regression Method: ", width=20)
l2.grid(row=1, column=0)
box2 = OptionMenu(window, variable2, "Linear Regression", "Decision Tree", "Polynomial
Regression", "KNN Regression",
"Random Forest")
box2.grid(row=1, column=1)
l2s = Label(window, text="Scaling Method: ", width=20)
l2s.grid(row=1, column=2)
box2s = Entry(window)
box2s.insert(END, "StandardScaler")
box2s.grid(row=1, column=3)

l2F = Label(window, text="Missing Values: ", width=20)
l2F.grid(row=1, column=4)
box2F = Entry(window)
box2F.insert(END, "drop")
box2F.grid(row=1, column=5)

variable3 = StringVar(window)
variable3.set("K-Means") # default value

l3 = Label(window, text="Clustering Method: ", width=20)
l3.grid(row=2, column=0)
box3 = OptionMenu(window, variable3, "K-Means")
box3.grid(row=2, column=1)
l3s = Label(window, text="Scaling Method: ", width=20)
l3s.grid(row=2, column=2)
box3s = Entry(window)
box3s.insert(END, "StandardScaler")
box3s.grid(row=2, column=3)

l3F = Label(window, text="Missing Values: ", width=20)
l3F.grid(row=2, column=4)
box3F = Entry(window, state='disabled')
box3F.insert(END, "drop")
box3F.grid(row=2, column=5)

l4 = Label(window, text="")
l4.grid(row=3, column=1, columnspan=2)

b1 = Button(window, text="Submit", width=40, fg="white", bg="blue",
command=buttonClick)
b1.grid(row=5, column=2, columnspan=2)

l6 = Label(window, text="")
l6.grid(row=6, column=1, columnspan=2)

l5 = Label(window, text="****Scaling Methods are: StandardScaler, MinMaxScaler,or
MaxScaler")
l5.grid(row=7, column=0, columnspan=8)

```

```
l5 = Label(window, text="***We deal with missing values by: drop, replaceMean,or  
replaceMode")  
l5.grid(row=8, column=0, columnspan=8)  
  
textl = Text(window, width=80, height=7)  
textl.grid(row=9, column=0, columnspan=30)  
  
# algorithm_classification = ""  
# algorithm_regression = ""  
# algorithm_cluster = ""  
# classificationScaler = ""  
# regressionScaler = ""  
# clusterScaler = ""  
  
window.mainloop()
```

Preprocessing.py

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import MaxAbsScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import LabelEncoder

class Preprocessor:
    def split(self, X, y, test_size):
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size,
random_state=58) # 58
        return X_train, X_test, y_train.astype(np.int64), y_test.astype(np.int64)

    def splitCluster(self, X, test_size):
        X_train, X_test = train_test_split(X, test_size=test_size)
        return X_train, X_test

    def scale(self, X_train, X_test, type):
        if type == "StandardScaler":
            standardScaler = StandardScaler()
            standardScaler.fit(X_train)
            X_train = standardScaler.transform(X_train)
            X_test = standardScaler.transform(X_test)
            return X_train, X_test

        elif type == "MinMaxScaler":
            minMaxScaler = MinMaxScaler()
            minMaxScaler.fit(X_train)
            X_train = minMaxScaler.transform(X_train)
            X_test = minMaxScaler.transform(X_test)
            return X_train, X_test

        elif type == "MaxScaler":
            maxScaler = MaxAbsScaler()
            maxScaler.fit(X_train)
            X_train = maxScaler.transform(X_train)
            X_test = maxScaler.transform(X_test)
            return X_train, X_test

        elif type == "RobustScaler":
            robustScaler = RobustScaler()
            robustScaler.fit(X_train)
            X_train = robustScaler.transform(X_train)
            X_test = robustScaler.transform(X_test)
            return X_train, X_test

    def encode(self, X, y, dataset_name):
        if dataset_name == 'diamonds':
            labelencoder = LabelEncoder()
            X[:, 1] = labelencoder.fit_transform(X[:, 1])
            X[:, 2] = labelencoder.fit_transform(X[:, 2])
            X[:, 3] = labelencoder.fit_transform(X[:, 3])
        elif dataset_name == 'cancer':
            labelencoder = LabelEncoder()
            y[:, 0] = labelencoder.fit_transform(y[:, 0])
        return X, y

    def drop(self, data):
```

```

data_modified = data.dropna(axis=0)
return data_modified

def replaceMean(self, data):
    data_modified = data.fillna(data.mean(), inplace=True)
    return data_modified

def replaceMode(self, data):
    data_modified = data.fillna(data.mode(), inplace=True)
    return data_modified

def dataCleaning(self, data, name):
    if name == 'cancer':
        data.drop(data.columns[0], axis='columns', inplace=True)
        X = data.iloc[:, 1:].to_numpy()
        y = data.iloc[:, 0].values.reshape(569, 1)

    elif name == 'diamonds':
        data.drop(data.columns[0], axis='columns', inplace=True)
        data_modified = data.drop(labels='price', axis=1)
        X = data_modified.iloc[:, 0:8].to_numpy()
        y = data_modified.iloc[:, 6].values.reshape(53940, 1)
    return X, y

```

regression.py

```
from sklearn import metrics
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import PolynomialFeatures

class Regression:
    def __init__(self, regressor_name, X, y):
        self.regressor_name = regressor_name
        self.X_train = X
        self.y_train = y

    def predict(self, X_test):
        if self.regressor_name == 'Decision Tree':
            decision_tree = DecisionTreeRegressor(random_state=0, max_depth=2)
            decision_tree.fit(self.X_train, self.y_train)
            return decision_tree.predict(X_test)

        elif self.regressor_name == 'Polynomial Regression':
            poly_features = PolynomialFeatures(degree=2)
            X_poly = poly_features.fit_transform(self.X_train)
            X_poly_test = poly_features.transform(X_test)
            poly_model = LinearRegression()
            poly_model.fit(X_poly, self.y_train)
            return poly_model.predict(X_poly_test)

        elif self.regressor_name == 'Linear Regression':
            linear_regressor = LinearRegression()
            linear_regressor.fit(self.X_train, self.y_train)
            return linear_regressor.predict(X_test)

        elif self.regressor_name == "Random Forest":
            # Random trees uses many random decision trees to output many different
            # we take the mean of the results
            random = RandomForestRegressor(n_estimators=100, max_depth=2)
            random.fit(self.X_train, self.y_train)
            return random.predict(X_test)

        elif self.regressor_name == 'KNN Regression':
            neigh = KNeighborsRegressor(n_neighbors=4)
            neigh.fit(self.X_train, self.y_train)
            return neigh.predict(X_test)

    def getScore(self, y_test, y_pred):
        if self.regressor_name == 'Polynomial Regression':
            return -metrics.r2_score(y_test, y_pred)
        else:
            return metrics.r2_score(y_test, y_pred)
```

classification.py

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics

class Classifier:
    def __init__(self, name, X, y):
        self.name = name
        self.X_train = X
        self.y_train = y
        self.fact = 0.32

    def classify(self, X_test):
        if self.name == 'KNN':
            n = KNeighborsClassifier(n_neighbors=4)
            n.fit(self.X_train, self.y_train)
            return n.predict(X_test)

        elif self.name == 'Decision Tree':
            decision_tree = DecisionTreeClassifier(max_depth=3)
            decision_tree.fit(self.X_train, self.y_train)
            return decision_tree.predict(X_test)

        elif self.name == 'Naive Bayes':
            b = GaussianNB()
            b.fit(self.X_train, self.y_train)
            return b.predict(X_test)

        elif self.name == "Random Forest":
            r = RandomForestClassifier(n_estimators=200, max_depth=3)
            r.fit(self.X_train, self.y_train)
            return r.predict(X_test)

    def calculateAccuracy(self, y_test, y_pred):
        return metrics.accuracy_score(y_test, y_pred) * 100
```

clustering.py

```
from sklearn.cluster import KMeans

class Clustering:
    def __init__(self, algorithm_cluster, X_train):
        self.algorithm_cluster = algorithm_cluster
        self.X_train = X_train

    def predict(self, X_test):
        kmeans = KMeans(n_clusters=3)
        kmeans.fit(self.X_train)
        return kmeans.predict(X_test)
```